

Data Transfer and Addressing

ELF memory maps, load/store architecture, addressing modes, and string/array traversal in AArch64

Topic 6 Outline

Topic Objective

Analysis of program memory organization (ELF layouts), AArch64 load-store architecture, memory addressing modes, alignment constraints, and pointer-based data traversal.

- **Section 1:** Program Memory Organization (ELF Layouts)
- **Section 2:** The Load-Store Architecture Paradigm
- **Section 3:** Load and Store Instructions (`ldr`, `str`)
- **Section 4:** AArch64 Memory Addressing Modes
- **Section 5:** Memory Access, Alignment, and Pair Transfers
- **Section 6:** Traversing Arrays, Pointers, and Strings
- **Section 7:** Summary and Worked Self-Test Exercises

Section 01

Program Memory Organization (ELF Layouts)

The ELF Memory Layout Paradigm

What is ELF?

ELF (Executable and Linkable Format) is the standard binary file format for executables, object code, and shared libraries in Unix-like systems.

- Defines how code, static data, and linking symbols are organized on disk and loaded into memory.
- Contains processor-specific flags (e.g., machine type identification like ARM architecture and alignment requirements).
- When an ELF binary executes, the operating system maps it into virtual memory segments.

The Linux Virtual Memory Map

- **Kernel Space:** Reserved space for kernel code, process tables, and kernel stacks.
- **User Stack:** Grows downward; stores local variables and activation frames.
- **Shared Libraries:** Mapped memory regions for dynamic libraries.
- **Runtime Heap:** Grows upward via dynamic allocation (`malloc`).
- **BSS / Data / Text:** Uninitialized, initialized, and executable code segments.

Process Virtual Memory

Kernel Virtual Memory

User Stack (↓)

Memory Mapped Region

Runtime Heap (↑)

Uninitialized Data (`.bss`)

Initialized Data (`.data`)

Program Text (`.text`)

The Load-Store Paradigm Architecture

Load-Store Architecture Principles

AArch64 follows a strict RISC load-store design philosophy:

- **Separation of Concerns:** Computation instructions (add, sub) process register operands exclusively. Memory instructions (ldr, str) handle data movement.
- **No In-Memory Arithmetic:** An instruction cannot simultaneously fetch from memory, add, and write back.
- **Simplified Datapath:** Keeps hardware control logic streamlined, highly predictable, and optimized for deep pipelining.

C Code Implication

C statements like `a = b + c;` require explicit loading of variables from memory into registers, addition, and a final store instruction.

Memory Addresses as Register Values

Memory locations are targeted using 64-bit virtual addresses held in general-purpose registers:

- The base address of a variable, array, or data structure is loaded into an X register (e.g., x1).
- Load/store instructions treat this register as a pointer referencing the target byte location.
- Pointer arithmetic can manipulate base addresses dynamically.

Pointer Register Rule

Only 64-bit general-purpose registers (x0–x30) can serve as base address pointers in AArch64.

Load and Store Instructions

Loading Data: The `ldr` Instruction

The Load Register (`ldr`) instruction copies data from memory into a target register:

- **Syntax Format:** `ldr destination_reg, [base_register]`
- **64-bit Transfer:** `ldr x0, [x1]` loads 8 bytes from the address in `x1` into `x0`.
- **32-bit Transfer:** `ldr w0, [x1]` loads 4 bytes into `w0`, zero-extending the upper 32 bits of `x0`.

Data Sizing Suffixes

Sub-word loads are supported via explicit qualifiers like `ldrb` (byte) and `ldrh` (halfword), with sign-extended variants (`ldrswb`, `ldrshw`).

Storing Data: The `str` Instruction

The Store Register (`str`) instruction copies data from a register into memory:

- **Syntax Format:** `str source_reg, [base_register]`
- **64-bit Transfer:** `str x0, [x1]` writes the 8-byte contents of `x0` to the memory address held in `x1`.
- **32-bit Transfer:** `str w0, [x1]` writes the lower 4 bytes of `w0` to memory.

Sub-Word Stores

Corresponding store variants `strb` (store byte) and `strh` (store halfword) allow granular writing of lower register bytes to memory destinations.

Precision Data Sizing and Sign Extension

AArch64 load instructions allow precise control over data width and extension behavior:

Instruction	Bytes Loaded	Extension Behavior
<code>ldrb w0, [x1]</code>	1 Byte (8 bits)	Zero-extended to 32 bits in <code>w0</code>
<code>ldrsb x0, [x1]</code>	1 Byte (8 bits)	Sign-extended to 64 bits in <code>x0</code>
<code>ldrh w0, [x1]</code>	2 Bytes (16 bits)	Zero-extended to 32 bits in <code>w0</code>
<code>ldrsh x0, [x1]</code>	2 Bytes (16 bits)	Sign-extended to 64 bits in <code>x0</code>

Sign Extension Importance

When loading signed 8-bit or 16-bit integers into larger registers, using sign-extended variants (`ldrsb/ldrsh`) preserves negative values correctly.

AArch64 Modes

Memory

Addressing

Addressing Mode Classification

An addressing mode defines how the CPU calculates the target memory address for a load or store instruction:

- **Base Register (Indirect):** Address is taken directly from the base register.
- **Immediate Offset:** Adds a constant offset to the base register without modifying the base.
- **Register Offset:** Adds a scaling or offset register to the base address.
- **Pre-Indexed:** Applies an offset, updates the base register, and accesses memory.
- **Post-Indexed:** Accesses memory using the base address, then updates the base register.

Base Register (Indirect) Addressing

Syntax and Semantics

```
ldr x0, [x1]
```

The memory address is taken strictly from the contents of the base register x1. The base register remains unchanged after execution.

- **Use Case:** Accessing a variable or dereferencing a pointer whose exact address is already loaded in a register.
- Analogous to pointer dereferencing in C: `*ptr`.

Immediate Offset Addressing

Syntax and Semantics

```
ldr x0, [x1, #16]
```

Computes the target address by adding an immediate constant offset (+#16) to the base register x1. The base register x1 is **not** modified.

- **Use Case:** Accessing fields within a structure or fixed offsets in a local stack frame.
- Immediate range rules depend on transfer size (must align with data width constraints).

Pre-Indexed Addressing Mode

Syntax and Semantics

```
ldr x0, [x1, #8] !
```

The immediate offset (+#8) is added to base register x1. This new calculated address is used for the memory load, **and** is written back into x1.

- Identified by the trailing exclamation mark (!).
- **Use Case:** Updating a pointer before accessing the next sequential memory item.

Post-Indexed Addressing Mode

Syntax and Semantics

```
ldr x0, [x1], #8
```

Data is loaded from the unmodified address in x1. After the memory access completes, the offset (+#8) is added to x1, updating the pointer register.

- Identified by placing the offset outside the square brackets.
- **Use Case:** Fetching current array elements and automatically advancing the pointer for the next iteration.

Register Offset and Scaled Addressing

AArch64 allows an index register to offset the base pointer, optionally scaled by data size:

- **Syntax:** `ldr x0, [x1, x2]` → Address is $x1 + x2$.
- **Scaled Register Offset:** `ldr x0, [x1, x2, lsl #3]` → Address is $x1 + (x2 \times 8)$.

Array Indexing Power

Scaling shifts the index register by the exact byte width of the data item ($2^3 = 8$ bytes for doublewords), mapping C array subscripts directly into hardware addressing.

Addressing Mode Summary

Mode	AArch64 Assembly	Address	Base Register
Offset	<code>ldr x0, [x1, x2]</code>	$x1 + x2$	Unchanged
Pre-index	<code>ldr x0, [x1, x2] !</code>	$x1 + x2$	$x1 = x1 + x2$
Post-index	<code>ldr x0, [x1], x2</code>	$x1$	$x1 = x1 + x2$

Note on Register vs. Immediate

Offset and index quantities can be specified as general-purpose registers (as shown above) or compatible immediate byte constants (e.g., #8).

Memory Access, Alignment, and Pair Transfers

Memory Alignment Rules in AArch64

Data alignment dictates how memory addresses map to data sizes:

- **Natural Alignment:** An n -byte data item must be stored at an address divisible by n .
 - Doublewords (8 bytes) require addresses divisible by 8.
 - Words (4 bytes) require addresses divisible by 4.
 - Halfwords (2 bytes) require addresses divisible by 2.
- **Strict Enforcement:** Unlike x86, unaligned general-purpose memory accesses in AArch64 can trigger hardware alignment fault exceptions.

Load and Store Pair Instructions (ldp / stp)

AArch64 provides specialized instructions to transfer **two registers simultaneously**:

- **Load Pair (ldp)**: `ldp x0, x1, [x2]` loads two 64-bit registers from consecutive memory locations starting at `x2`.
- **Store Pair (stp)**: `stp x0, x1, [x2]` stores two 64-bit registers consecutively to memory.

Performance Efficiency

Pair instructions utilize wide memory data buses, moving 16 bytes in a single instruction cycle and reducing overall instruction count.

Stack Operations via Pair Transfers

Stack management relies heavily on pre-index and post-index pair instructions:

Pushing to the Stack (Pre-Index)

```
stp x29, x30, [sp, #-16]!
```

Decrements stack pointer (sp) by 16 bytes, then stores frame pointer (x29) and link register (x30).

Popping from the Stack (Post-Index)

```
ldp x29, x30, [sp], #16
```

Loads registers from current sp, then increments stack pointer by 16 bytes, restoring stack balance.

Traversing Arrays, Pointers, and Strings

Traversing Arrays, Pointers, and Strings

Initializing Base Pointers in C and Assembly

C Pointer Initialization

- `int array[6] = {2, 3, 11, 4, 55, 6};`
- `int *ptr = array; // base pointer`
- `*(ptr += 4) += 10;`

AArch64 Base Register Setup

- Declare static data arrays in the `.data` section.
- Load the array base address into a pointer register using symbol directives: `ldr x0, =array`

Mapping C Arrays to AArch64 Assembly

Arrays consist of contiguous memory locations holding elements of identical size:

- Let `x1` hold the base address of an integer array (`int A[]`, where each element is 4 bytes).
- Accessing element `A[i]` requires computing the effective address:
 $\text{Base} + (\text{Index} \times \text{Element Size})$.
- Assembly implementations utilize scaled register offsets or pointer arithmetic loops.

Example Array Access

`ldr w0, [x1, x2, lsl #2]` loads `A[i]` into `w0`, where `x2` holds index `i` scaled by $2^2 = 4$.

Iterating Over Strings (.asciz)

Strings in C/C++ are represented as contiguous character arrays terminated by a null byte (0):

- Declared using assembler directives: `message: .asciz "Hello World!"`
- Accessing individual characters requires byte-level instructions like `ldrb` (Load Register Byte).
- Using post-indexed byte loads (`ldrb w1, [x0], #1`) fetches the current character and automatically advances the pointer in one instruction.

Loop Termination Check

The loop continues reading until the loaded character equals zero (`cmp w1, #0`, followed by `b.eq end`).

Mapping C Arrays to AArch64 Assembly

Arrays consist of contiguous memory locations holding elements of identical size:

- Let `x1` hold the base address of an integer array (`int A[]`, where each element is 4 bytes).
- Accessing element `A[i]` requires computing the effective address:
 $\text{Base} + (\text{Index} \times \text{Element Size})$.
- Assembly implementations utilize scaled register offsets or pointer arithmetic loops.

Example Array Access

`ldr w0, [x1, x2, lsl #2]` loads `A[i]` into `w0`, where `x2` holds index `i` scaled by $2^2 = 4$.

Iterating Over Strings (.asciz)

Strings in C/C++ are represented as contiguous character arrays terminated by a null byte (0):

- Declared using assembler directives: `message: .asciz "Hello World!"`
- Accessing individual characters requires byte-level instructions like `ldrb` (Load Register Byte).
- Using post-indexed byte loads (`ldrb w1, [x0], #1`) fetches the current character and automatically advances the pointer in one instruction.

Loop Termination Check

The loop continues reading until the loaded character equals zero (`cmp w1, #0`, followed by `b.eq end`).

Summary & Self-Test

Topic 6 Comprehensive Summary

- **ELF Memory Layout:** Programs execute across mapped virtual segments including Text, Data, BSS, Heap, and Stack.
- **Load-Store Paradigm:** Memory access is restricted to explicit `ldr` and `str` instructions; computation happens in registers.
- **Addressing Modes:** Supports base indirect, immediate offset, register offset, pre-indexed, and post-indexed pointer manipulation.
- **Alignment & Pairs:** Natural alignment is strictly enforced; `ldp` and `stp` optimize stack frames.
- **Data Structures:** Scaled offsets, pointers, and null-terminated byte loops (`ldrb`) govern array and string traversal.

Self-Test: Questions 1 & 2

Question 1: ELF Segments

What is the primary difference between the `.text` segment and the `.data` segment in an ELF executable?

Question 2: Immediate vs. Indexed Addressing

What is the primary operational difference between immediate offset addressing (`[x1, #8]`) and pre-indexed addressing (`[x1, #8]!`)?

Self-Test: Solutions 1 & 2

Solution 1

The `.text` segment stores the program's executable machine instructions (read-only), while the `.data` segment stores initialized global and static variables (read/write).

Solution 2

Immediate offset adds the constant to the base register solely to calculate the target address, leaving the base register unchanged. Pre-indexed addressing calculates the address **and writes the updated address back** into the base register.

Self-Test: Questions 3 & 4

Question 3: Post-Indexed String Traversal

How does the instruction `ldrb w1, [x0], #1` simplify iterating through a null-terminated string compared to using a separate addition instruction?

Question 4: Memory Alignment

What happens when an unaligned 64-bit doubleword load is executed on standard AArch64 hardware?

Self-Test: Solutions 3 & 4

Solution 3

It loads the byte at the current address in `x0` into `w1` and automatically increments `x0` by 1 in a single instruction cycle, eliminating the need for a separate `add x0, x0, #1` statement.

Solution 4

Unaligned accesses violate natural alignment rules and typically trigger a hardware alignment fault exception, crashing the program or invoking an OS trap handler.

Self-Test: Questions 5 & 6

Question 5: Pair Transfers on the Stack

How do `stp` and `ldp` manage stack pointer updates during function entry and exit?

Question 6: Scaled Register Addressing

In the instruction `ldr x0, [x1, x2, lsl #3]`, what role does the `lsl #3` modifier play in indexing an array of 64-bit integers?

Self-Test: Solutions 5 & 6

Solution 5

Using pre-indexed store pair (`stp ... [sp, #-16]!`), the stack pointer is decremented and registers are saved in one atomic step. Post-indexed load pair (`ldp ... [sp], #16`) restores registers and re-increments the stack pointer.

Solution 6

The `lsl #3` scales the index in `x2` by multiplying it by $2^3 = 8$ (the byte size of a 64-bit doubleword), correctly translating logical array indices into byte offsets.

Review: Addressing Mode Summary Table

Addressing Mode	Assembly Syntax Example	Base Register State After Access
Base Register	<code>ldr x0, [x1]</code>	Unchanged
Immediate Offset	<code>ldr x0, [x1, #16]</code>	Unchanged
Pre-Indexed	<code>ldr x0, [x1, #8] !</code>	Updated ($\text{Base} + 8$)
Post-Indexed	<code>ldr x0, [x1], #8</code>	Updated ($\text{Base} + 8$)

Wrap-Up & Transition to Topic 7

Moving Forward to Topic 7

Now that we have mastered memory data transfer, ELF layouts, addressing modes, and string/array traversal, our next session will explore:

- **Control Flow & Branching:** Unconditional and conditional branches in AArch64.
- **Subroutine Linkage:** Function calls (`bl`), return linkage (`ret`), and ABI parameter passing conventions.

End of Topic 6 — Questions & Discussion

Topic 6 Source Notes

Source Reference Verification

- ELF memory mapping and program segmentation conform to standard System V ABI and Linux specifications.
- Load-store mechanics, string handling (`ldr`), and offset calculations follow standard ARMv8/ARMv9 AArch64 ISA guidelines.